

Myntra Visual Search

Decision Summary — Problem to Recommendation, in Numbers

Sambhavya Nayak • github.com/SambhavyaNayak/Myntra-Visual-Search-Case-Study

The problem, in one line

Myntra's visual search treats a whole photo as one query, so photos with more than one garment return blended, low-relevance results — even though image search is one of Myntra's fastest-growing entry points, up **35% year-over-year**.

The decision

Ship tap-to-select visual search: auto-detect individual garments in a photo, let the user tap the one they want, then refine with attribute chips (price, color, fabric) — a pattern proven at scale by Pinterest Lens, adapted for fashion-specific needs.

Why this option, backed by the numbers

Option	Reach/qtr	Impact	Confidence	Effort	RICE Score
A: Tap-to-Select (chosen)	2,000,000	2 (High)	70%	4 mo	700,000
B: Manual crop only	2,000,000	1 (Med)	80%	1.5 mo	1,066,667
C: Filter chips only	5,000,000	0.5 (Low)	85%	1 mo	2,125,000

The numbers alone point to Option C — but that's the wrong read. C scores highest only because it's cheap and touches more users, not because it solves the problem: it filters an already-blended result set instead of fixing the ambiguity causing bad results in the first place. That's the core judgment call in this case study: a lower RICE score was chosen deliberately because Impact was scored honestly (0.5 for C vs. 2 for A), and a real product decision weighs *why* a score is low, not just its rank.

Final call: Option A (Tap-to-Select), with Option C's filter chips folded in as a post-selection refinement step, and Option B's manual crop kept as a fallback inside Option A rather than shipped alone.

What proves this was the right call, post-launch

- **Primary metric:** visual-search-to-cart conversion, measured only on sessions with 2+ detected garments.
- **Guardrails:** search latency (target: no more than +300ms added) and single-item-photo conversion (should stay flat).
- **Validation:** A/B test, user-randomized, with a Sample Ratio Mismatch (SRM) check before trusting any result — same integrity check used in the Cookie Cats case study.

Assumptions used (transparency note)

The reach, effort, and confidence figures above are case-study estimates, not real Myntra data. They're built on one verifiable public fact (35% YoY growth in image search) and reasonable assumptions from there — used to demonstrate how the RICE methodology and trade-off reasoning work, not to claim real company metrics.

What I'd do differently with more time or real data

- Validate the “2M multi-item searches/quarter” estimate against real search logs before committing engineering effort — this number drives the entire prioritization and would be the first thing to fact-check on a real team.
- Run a small qualitative test (5-10 users) on a prototype before the full A/B test, to catch obvious usability issues in the tap-to-select interaction cheaply, before the full 4-month build.